

# CS1020E Tutorial + Lab 10

Mark NG

`a0116298@u.nus.edu`  
`http://mollymr305.github.io`

November 04, 2016

# Tutorial Solutions

## “Tutorial 10 – Hashing”

First see

<https://visualgo.net/hashtable>

## Question 1: Simulation

In this question we will simulate the operations `insert(key)` and `erase(key)` on a `std::unordered_set`, denoted by the shorthand  $I(k)$  and  $D(k)$  respectively. Note that a **Hash Table**, e.g. `unordered_map`, works on a **<Key, Value> pair**, while in a **Hash Set**, or `unordered_set`, the **value is the key itself**.

Fill the contents of the hash table after each insert / delete operation.

The Hash Table has “table size” of 5, i.e. 5 buckets. The hash function is  $h(key) = key \% 5$

## Question 1 (a): Problem

Use linear probing as the collision resolution technique:

	<b>0</b>	<b>1</b>	<b>2</b>	<b>3</b>	<b>4</b>
I(7)					
I(12)					
I(22)					
<b>D(12)</b>					
I(8)					

## Question 1 (a): Solution

Use linear probing as the collision resolution technique:

	<b>0</b>	<b>1</b>	<b>2</b>	<b>3</b>	<b>4</b>
I(7)			<b>7</b>		
I(12)			7	<b>12</b>	
I(22)			7	12	<b>22</b>
<b>D(12)</b>			7	<del>12</del>	<b>22</b>
I(8)			7	<b>8</b>	22

## Question 1 (b): Problem

Use quadratic probing as the collision resolution technique:

	<b>0</b>	<b>1</b>	<b>2</b>	<b>3</b>	<b>4</b>
I(7)					
I(12)					
I(22)					
<b>I(2)</b>					

## Question 1 (b): Solution

Use quadratic probing as the collision resolution technique:

	<b>0</b>	<b>1</b>	<b>2</b>	<b>3</b>	<b>4</b>
<b>I(7)</b>			<b>7</b>		
<b>I(12)</b>			7	<b>12</b>	
<b>I(22)</b>		<b>22</b>	7	12	
<b>I(2)</b>	Unable to find free slot				



## Question 1 (c): Problem

Use double hashing as the collision resolution technique,  $h_2(\text{key}) = \text{key} \% 3$ :

	<b>0</b>	<b>1</b>	<b>2</b>	<b>3</b>	<b>4</b>
I(7)					
I(22)					
I(12)					

## Question 1 (c): Solution

Use double hashing as the collision resolution technique,  $h_2(\text{key}) = \text{key} \% 3$ :

	<b>0</b>	<b>1</b>	<b>2</b>	<b>3</b>	<b>4</b>
I(7)			<b>7</b>		
I(22)			7	<b>22</b>	
<b>I(12)</b>	Infinite loop on collision resolution as $h_2(12) == 0$				

## Question 1 (d): Problem

Use double hashing as the collision resolution technique,  $h_2(\text{key}) = 7 - (\text{key} \% 7)$ :

	<b>0</b>	<b>1</b>	<b>2</b>	<b>3</b>	<b>4</b>
<b>I(7)</b>					
<b>I(12)</b>					
<b>I(22)</b>					
<b>I(2)</b>					

## Question 1 (d): Solution

Use double hashing as the collision resolution technique,  $h_2(\text{key}) = 7 - (\text{key} \% 7)$ :

	<b>0</b>	<b>1</b>	<b>2</b>	<b>3</b>	<b>4</b>
<b>l(7)</b>			<b>7</b>		
<b>l(12)</b>			7		<b>12</b>
<b>l(22)</b>			7	<b>22</b>	12
<b>l(2)</b>	Infinite loop on collision resolution as $h_2(2) \% m == 0$				

## Question 2: Hashing or No Hashing

Hash Table is an ADT that allows for find, insert, and delete operations in  $O(1)$  average-case time, if properly designed. However, it is not without its limitations. For each of the cases described below, state if Hash Table can be used.

In each case, if it is possible to use Hash Table, describe its design, including:

- 1) The <Key, Value> pair
- 2) Hashing function/algorithm
- 3) Collision resolution (including second hash function/algorithm for double hashing)

Otherwise, why is Hash Table not suitable for that particular case?

## Question 2 (a):

**(a)** A bank has many bank accounts. Each account has name, account number, and current balance. The operations to perform are: retrieve balance, deposit money, and withdraw money. For each operation, either the name or account number will be given. The amount will also be given for the last two operations.

## Question 2 (a):

**(a)** A bank has many bank accounts. Each account has name, account number, and current balance. The operations to perform are: retrieve balance, deposit money, and withdraw money. For each operation, either the name or account number will be given. The amount will also be given for the last two operations.

**(a)** Yes. 2 Hash Tables for efficient lookup using either key

<Key, Value> Pair: <name, BankAccount\*> and <account number, BankAccount\*>

Hashing Algorithm: string hashing algorithm  $f(\text{key}, \text{charWeight}, \text{tableSize})$ , that takes into account every character weighted by its position, modulo the size of the hash table

Collision Resolution: Double hashing using  $h_2(\text{name}) = 1 + f_2(\text{name}, \text{weight}, m)$ , where  $f_2$  has a different character weighting from  $f$ , and the modulo  $m$  used is smaller than table size

We are here assuming that the table size is a prime number, so that every  $h_2(\text{name})$  is co-prime to it.

## Question 2 (b):

**(b)** A mini-population census is to be conducted on every person in your neighbourhood. We are only interested in storing every person's name and age. The operations to perform are: retrieve age by name, and retrieve names by age.



## Question 2 (b):

**(b)** A mini-population census is to be conducted on every person in your neighbourhood. We are only interested in storing every person's name and age. The operations to perform are: retrieve age by name, and retrieve names by age.

**(b)** Yes, using 2 Hash Tables for efficient lookup both ways.

<Key, Value> Pair: <name, age> and <age, names>, to find age and names respectively

Hashing Algorithm:  $h(\text{name})$  = standard string hashing, and,  $h(\text{age})$  = age (direct addressing)

Collision Resolution:

<name, age> Hash Table – Double hashing using  $h_2(\text{name})$  as in part (a)

<age, name> Hash Table – Separate chaining for names with the same age

## Question 2 (c):

**(c)** A larger population census is also conducted across the country, similarly containing only every person's name and age. The operation to perform is: retrieve names of people eligible for voting. Only people above a certain age are eligible for voting. However, we still need to store the rest of the data.

## Question 2 (c):

(c) A larger population census is also conducted across the country, similarly containing only every person's name and age. The operation to perform is: retrieve names of people eligible for voting. Only people above a certain age are eligible for voting. However, we still need to store the rest of the data.

(c) Yes, **direct addressing** table as in part (b)'s <age, names> Hash Table. The operation can be performed by iterating through age, traversing the chain/list for each age, starting from the minimum eligible age.

## Question 2 (d):

**(d)** A different population census similarly contains only the name and age of every person. The operation to perform is: retrieve people with a given last name.

## Question 2 (d):

**(d)** A different population census similarly contains only the name and age of every person. The operation to perform is: retrieve people with a given last name.

**(d)** Yes, if we can extract the last name from a name.

<Key, Value> Pair: <last name, pair<full name, age>s>

Hashing Algorithm:  $h(\text{name}) = \text{standard string hashing}$

Collision Resolution: Separate chaining for people with the same last name

## Question 2 (e):

**(e)** In a scientific experiment, data on energy converted for different runs is collected. Each run has a different velocity and distance, both floating point numbers. The operation to perform is: retrieve the energy value for a given velocity and distance. [Hint: What would happen when you hash a floating-point number?]

## Question 2 (e):

**(e)** In a scientific experiment, data on energy converted for different runs is collected. Each run has a different velocity and distance, both floating point numbers. The operation to perform is: retrieve the energy value for a given velocity and distance. [Hint: What would happen when you hash a floating-point number?]

**(e)** Not a good choice, as floating-point numbers are an **approximation** of real numbers. For example, 0.1 cannot be represented as a binary floating point number. It would be difficult to design a good hash function that evenly distributes all keys in general.

However, if we allow precision of up to a number of binary digits, then this may be possible. We can then map a key to a particular address, and find a value that is approximately equal to that key. We will probably have to write our own hash function, as a function of both numbers.

## Question 3: Hash Function

A good hash function is essential for good hash table performance. A good hash function is easy/efficient to compute and will evenly distribute the possible keys. Comment on the flaw (if any) of the following hash functions. Assume the load factor  $\alpha$  = number of keys / table size = 0.3 for all the following cases:



## Question 3 (a):

- (a) The hash table has size 100. The keys are positive even integers. The hash function is  
 **$h(key) = key \% 100$**

## Question 3 (a):

- (a) The hash table has size 100. The keys are positive even integers. The hash function is

$$h(\text{key}) = \text{key} \% 100$$

- (a) No key will be hashed directly to **odd-numbered slots** in the table, resulting in wasted space, and a higher number of collisions in the remaining slots than necessary.

Aside from the hash function, the hash table size may not be good as it is not a prime number. If there are many keys that have identical last two digits, then they will all be hashed to the same slot, resulting in many collisions.

## Question 3 (b):

(b) The hash table has size 49. The keys are positive integers. The hash function is

$$h(key) = (key * 7) \% 49$$

## Question 3 (b):

(b) The hash table has size 49. The keys are positive integers. The hash function is

$$h(key) = (key * 7) \% 49$$

(b) All keys will be hashed only into slots 0, 7, 14, 21, 28, 35 and 42. Also, as in (a), the hash table size is not a prime number.

## Question 3 (c):

- (c) The hash table has size 100. The keys are non-negative integers in the range of  $[0, 10000]$ . The hash function is

$$h(key) = \lfloor \sqrt{key} \rfloor \% 100$$

## Question 3 (c):

- (c) The hash table has size 100. The keys are non-negative integers in the range of  $[0, 10000]$ . The hash function is

$$h(key) = \lfloor \sqrt{key} \rfloor \% 100$$

- (c) Keys are **not uniformly** distributed. Many more keys are mapped to the larger-indexed slots. Also, as in (a), the hash table size is not a prime number.

## Question 3 (d):

(d) The hash table has size 1009. The keys are valid email addresses. The hash function is

**$h(key) = (\text{sum of ASCII values of each of the last 10 characters}) \% 1009$**

See <http://www.asciitable.com> for ASCII values

## Question 3 (d):

(d) The hash table has size 1009. The keys are valid email addresses. The hash function is

**$h(\text{key}) = (\text{sum of ASCII values of each of the last 10 characters}) \% 1009$**

See <http://www.asciitable.com> for ASCII values

(d) Keys are **not evenly distributed** because many email addresses have the **same domain names**, e.g. "u.nus.edu", "gmail.com". Many email addresses will be hashed to the same slot, resulting in many collisions.



## Question 3 (e):

(e) The hash table has size 101. The keys are integers in the range of  $[0, 1000]$ . The hash function is

$$h(\text{key}) = \lfloor \text{key} * \text{random} \rfloor \% 101, \text{ where } 0.0 \leq \text{random} \leq 1.0$$

## Question 3 (e):

- (e) The hash table has size 101. The keys are integers in the range of  $[0, 1000]$ . The hash function is  $h(key) = \lfloor key * random \rfloor \% 101$ , where  $0.0 \leq random \leq 1.0$
- (e) This hash function is **not deterministic**. The hash function does not work because, while using a **given key** to retrieve an element after inserting it into the hash table, we **cannot always reproduce the same address**.

## Question 4: Problem

You have a group of non-empty strings, with each string having an integer ID associated with it. The strings are stored in `unordered_map<int, string> map`. Which of these 3 ways successfully finds whether an element with key 3 exists in `map`?

Read the documentation and try it out. Remember to compile using the C++11 standard.

```
bool found = (map.find(3) != map.end());
```

```
bool found = (map.at(3) != "");
```

```
bool found = (map[3] != "");
```

How about doing the same in an `unordered_set<int>`?

## Question 4: Solution

If we do not know whether an element with a given key exists, we should only use `find()`.

`find()` returns an iterator 'pointing' to the key-value pair if found, or an iterator equivalent to `end()` if no such key is found.

When no such key exists:

- `operator[]` creates a new element, "" in this case, adding an additional element to the map
- `at()` throws an Exception, a runtime error that may stop the program

Both `unordered_map<K, V>` and `unordered_set<K>` have `insert()`, `erase()`, and `find()`.  
`unordered_set<K>`, however, does not have `at()` or `operator[]`.

# End of Tutorial Discussion

**Note:** Detailed solutions (i.e. the file T10\_ans.pdf) will be released soon at

<http://www.comp.nus.edu.sg/~stevenha/cs1020e.html>

# Take Home Lab

Use hash map...

Let's take a short break!

# Practical Exam Practice



# Practical Exam Practice

- Write a function which takes in a positive integer and returns the *sum* of all its digits. Solve this recursively.

# Practical Exam Practice

- Write a function which takes in a positive integer and returns the *sum* of all its digits. Solve this recursively.
- Write a function which takes in a positive integer and returns the *product* of all its digits. Solve this recursively.

# Kattis Problem

Solve this problem **as fast as possible**

# Kattis Problem

Solve this problem **as fast as possible**

`https://open.kattis.com/problems/fizzbuzz`

# Kattis Problem

Solve this problem.

# Kattis Problem

Solve this problem.

<https://open.kattis.com/problems/busyschedule>

# Any Questions?

Good luck for your Practical Exam. See you next week!